

# cse5441 - parallel computing

## Open MP

# OpenMP -vs- pthreads

- both use **global shared memory** paradigm
- both require rigorous analysis of **variable scope and sharing**
- both methods operate at the “**thread**” level
- pthreads design objective: to give programmer **maximal control**
- OpenMP design objective: to provide semi-automatic parallelization **without change to existing code**
- **performance is highly variable**, and implementation specific

context:

**global variables** – with respect to a parallel region

# what is OpenMP?

essentially:

- a set of compiler directives
- a set of library routines
- aims to implement **best practices** from 20 years of threads experience

available for C, C++ and Fortran

# OpenMP core syntax

- Most of the constructs in OpenMP are compiler directives:

**#pragma omp construct [clause [clause]...]**

- Example

**#pragma omp parallel num\_threads(4)**

- Function prototypes and types in the file: **#include <omp.h>**

- Most OpenMP constructs apply to a “**structured block**”.

- **Structured block:** a block of one or more statements surrounded by “{ }”, with one point of entry at the top and one point of exit at the bottom.

(It's OK to have an **exit()** within the structured block . . .

. . .

but all will exit.)

# hello world ☺

```
#include <omp.h>
void main()
{
    #pragma omp parallel
    {
        int ID = 0;
        printf(" hello(%d) ", ID);
        printf(" world(%d) \n", ID);
    }
}
```

compile openMP programs on stdlinux with **-fopenmp**


# hello world 😊

```
#include <omp.h>
void main()
{
    #pragma omp parallel
    {
        int ID = omp_get_thread_num();
        printf(" hello(%d) ", ID);
        printf(" world(%d) \n", ID);
    }
}
```

# thread interaction

- multi-threaded global shared memory model
- thread communication by shared variables (shared memory)
- unintended sharing causes race conditions
  - inconsistent and unpredictable output
  - avoided with synchronization and serial segments

# fork – join parallelism



note: parallel regions “can” be nested:

[https://docs.oracle.com/cd/E19059-01/stud.10/819-0501/2\\_nested.html](https://docs.oracle.com/cd/E19059-01/stud.10/819-0501/2_nested.html)



# parallel region data distribution

“workload distribution”

```
double A[1000];

omp_set_num_threads(4);
#pragma omp parallel
{
    int t_id = omp_get_thread_num();
    for (int i = t_id; i < 1000; i += omp_num_threads())
    {
        A[i] = 0;
    }
}
```

**cyclic  
distribution**

```
double A[1000];

omp_set_num_threads(4);
#pragma omp parallel
{
    int t_id = omp_get_thread_num();
    int b_size = 1000 / omp_num_threads();
    for (int i = t_id * b_size; i < (t_id+1) * b_size; i++)
    {
        A[i] = 0;
    }
}
```

**block  
distribution**

we will be  
*finessing*  
*boundary*  
*conditions ...*

# OpenMP requesting threads

```
double A[1000];

#pragma omp parallel num_threads(my_choice)
{
    int t_id = omp_get_thread_num();
    for (int i = t_id; i < 1000; i += omp_num_threads())
    {
        A[i] = 0;
    }
}
```

```
{
    each thread will
    execute the code
    within the block
}
```

implicit barrier




# Tim Mattson & Larry Meadows

- Principal Engineers for Intel
- active part of OpenMP architecture review board
- provided content on many slides
- <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>
- (also on Carmen ...)

the software professional must design for  
a specific OpenMP version

# numerical integration

background



Mathematically:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

Which can be approximated by:

$$\sum_{i=0}^n H(x_i) \Delta x \approx \pi$$

where each rectangle has width  $\Delta x$  and height  $H(x_i)$  at the middle of interval  $i$ .

# serial pi

example 1

```
int      num_steps = 100000;
double   step;

void main ()
{
    int      i;
    double   x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;
    for (i = 0; i < num_steps; i++)
    {
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```



# SPMD pattern with OpenMP

- **Single Program Multiple Data**
- each thread runs same program
  - selection of data, or braching conditions, based on `thread id`
- in OpenMP implementations:
  - parallelize loops
  - query `thread_id` and `num_threads`
  - `partition` input data among threads

“workload distribution”

# it's your turn ... parallel pi

exercise 1

using:

- `#pragma omp parallel`
- `(void) omp_set_num_threads(int)`
- `(int) omp_get_num_threads()`
- `(int) omp_get_thread_num()`

create a parallel version of pi  
by parallelizing the *for* loop

pay close attention to variables which will be shared

# parallel assignment race condition

**sum = sum + 4.0/(1.0+x\*x);**

load_register	1, @sum
set_register	2, 4.0
set_register	3, 1.0
load_register	4, @x
multiply_into	5, 4, 4
add_into	4, 3, 5
divide_into	3, 2, 4
add_into	2, 1, 3
store	2, @sum



# parallel pi

exercise 1

```
#include <omp.h>
int      num_steps = 100000;
double   step;
#define   NUM_THREADS      2

void main ()
{
    int      i, nthreads;
    double   pi = 0.0, sum[NUM_THREADS];

    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
```

another way to request ^  
a number of threads

this loop is serial ->

```
#pragma omp parallel
```

```
{
    int      i, id, nthrds;
    double   x;
```

```
    id      = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthrds = nthrds;
```

<- example  
promotion

```
    sum[ id ] = 0.0;
    for ( i = id; i < num_steps; i += nthrds)
```

```
    {
        x = (i+0.5)*step;
        sum[id] += 4.0/(1.0+x*x);
    }
```

^ partition method  
cyclic distribution

```
} <- implicit barrier
```

```
for( i = 0; i < nthreads; i++)
```

```
{
    pi += sum[i] * step;
```

a "reduction"

```
}
```

# parallel updates – cache considerations

```
sum[id] += 4.0/(1.0+x*x);
```

```
sum[id] = sum[id] + 4.0/(1.0+x*x);
```

## false sharing

when updating:

- consider multi-processors with local caches
- consecutive elements assigned to different threads will likely go to other caches
- write will update local cache, but all other local caches would be marked invalid

# cyclic -vs- block distribution

```
double A[1000];

omp_set_num_threads(4);
#pragma omp parallel
{
    int t_id = omp_get_thread_num();
    for (int i = t_id; i < 1000; i += omp_num_threads())
    {
        sum[id] += 4.0/(1.0+x*x);
    }
}
```

```
double A[1000];

omp_set_num_threads(4);
#pragma omp parallel
{
    int t_id = omp_get_thread_num();
    int b_size = 1000 / omp_num_threads();
    for (int i = (t_id-1) * b_size; i < t_id * b_size; i++)
    {
        sum[id] += 4.0/(1.0+x*x);
    }
}
```

# critical code regions

```
#pragma omp parallel
{
float    B;
int      i, id, nthrds;

    id      = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    for( i = id; i < MAX; i + nthrds)
    {
        B = big_job(i);
        #pragma omp critical
        {
            consume (B, res);
        }
    }
}
```

# it's your turn ... parallel pi

exercise 2

using:

- `#pragma omp parallel`
- `(void) omp_set_num_threads(int)`
- `(int) omp_get_num_threads()`
- `(int) omp_get_thread_num()`
- `#pragma omp critical`
- `block workload distribution`

create a parallel version of pi which:

- does not exhibit false sharing
- does not use an array of sums

# parallel pi - no false sharing

exercise 2

```
int      num_steps = 100000;
double   step;
#define  NUM_THREADS      2

void main ()
{
    int      i, nthreads;
    double   pi = 0.0;

    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
```

no array, no false sharing ->

```
#pragma omp parallel
{
    int      i, id, nthrds;
    double   x, sum;  <- sum is now local

    id      = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthreads = nthrds;

    sum = 0.0;
    for ( i = id; i < num_steps; i += nthrds)
    {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
    #pragma omp critical
    {
        pi += sum * step;
    }
}
```

^ each thread adds its partial  
sum one thread at a time

# named critical code regions

```
#pragma omp parallel
{
float    B;
int      i, id, nthrds;

    id      = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    for( i = id; i < MAX; i + nthrds)
    {
        B = big_job(i);
        #pragma omp critical consume_it
        {
            consume_1 (B, res);
label 1: consume_2(B, upd);
            consume_3(res, upd);
        }
    }
}
```

goto label1;

branching into  
critical region  
prohibited



# atomic code regions

```
int      x;

#pragma omp parallel
{
    int      i, id, nthrds, y;

    id      = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    for( i = id; i < MAX; i + nthrds)
    {
        y = something(fun);
        /* Protect against race conditions */
        #pragma omp atomic
        {
            x += y;
        }
    }
}
```



# cse5441 - parallel computing

## Open MP

# OpenMP loop worksharing

```
#pragma omp parallel
{
    // id      = omp_get_thread_num();
    // nthrds = omp_get_num_threads();
    // for( i = id; i < MAX; i + nthrds)
    #pragma omp for
    for( i = 0; i < MAX; i++)
    {
        B = big_job(i);
    }
}
```

# combined parallel / worksharing construct

```
#pragma omp parallel for
{
    for( i = 0; i < MAX; i++)
    {
        B = big_job(i);
    }
}
```

# OpenMP reduction

How do we handle this case?:

```
double  avg = 0.0;
double  A[SIZE];
#pragma omp parallel for
for (int i = 0; i < SIZE; i++)
{
    avg += A[ i ];
}

avg = avg / SIZE;
```

We can use the clause:

`reduction(op : list)`

# OpenMP reduction

```
double  avg = 0.0;
double  A[SIZE];
#pragma omp parallel for reduction( + : avg )
for (int i = 0; i < SIZE; i++ )
{
    avg += A[ i ];
}

avg = avg / SIZE;
```

parallel region

- each variable in *list* is made thread local
- each *list* variable is initialized appropriate to *op*
- local *list* variables combined according to *op* at end of loop
- combined value exported to original “global” variable

# it's your turn ... parallel pi

exercise 3

- **Example**

```
#pragma omp parallel num_threads(4)
```

using:

- *#pragma omp parallel for*
- *reduction ( op : list )*
- *num\_threads ( int )*
- 8 threads

create a parallel version of pi which  
is as close to serial pi as possible

# parallel pi

exercise 3

## serial pi

```
int      num_steps = 100000;
double   step;
```

```
void main ()
```

```
{
  int      i;
  double   x, pi, sum = 0.0;
```

```
    step = 1.0/(double) num_steps;
```

parallelize, and reduce into sum →

```
    for (i = 0; i < num_steps; i++)
```

```
    {
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
```

```
    pi = step * sum;
```

```
}
```

## parallel pi

```
int      num_steps = 100000;
double   step;
```

```
void main ()
```

```
{
  int      i;
  double   x, pi, sum = 0.0;
```

```
    step = 1.0/(double) num_steps;
```

```
    #pragma omp parallel for private(x) reduction( + : sum) \
        num_threads(8)
```

```
    for (i = 0; i < num_steps; i++)
```

```
    {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
```

```
    pi += sum * step;
```

```
}
```

line  
continuation

v

# cse5441 - parallel computing

Open MP

additional  
topics



# setting num\_threads at run time

```
$g++ -fopenmp my_omp_program.cc  
$ a.out 4 <my_data_file
```

```
...
```

```
int main( int argc, char* argv[ ] )  
{  
    nt = atoi( argv[ 1 ] );
```

```
...
```

```
#pragma omp parallel num_threads(nt)  
{
```

```
...
```

<-  
pragma supports  
variables

# OpenMP synchronization: barriers

```
#pragma omp parallel private(id)
{
    id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier
```

explicit barrier →

```
#pragma omp for
for(i=0; i<N; i++)
{
    C[i]=big_calc3(i,A);
}
```

implicit barrier at end  
of parallel region →

```
#pragma omp for nowait
for(i=0; i<N; i++)
{
    B[i]=big_calc2(C, i);
}
A[id] = big_calc4(id);
}
```

no barrier!  
**nowait** cancels barrier creation →

NOTE:  
no implicit barriers at the  
beginning of pragmas

what could happen following  
nowait clause?

# OpenMP synchronization: master

```
#pragma omp parallel  
{
```

multiple threads of control →

```
do_many_things();
```

```
#pragma omp master
```

only **master thread** →  
executes this region

```
{  
    reset_boundaries();  
} ← no implicit barrier here
```

multiple threads of control →

```
do_many_other_things();
```

```
}
```

# OpenMP synchronization: single

```
#pragma omp parallel  
{
```

multiple threads of control →

```
do_many_things();
```

```
#pragma omp single
```

a single thread is chosen  
to execute this region →

```
{  
    reset_boundaries();  
} ← implicit barrier
```

multiple threads of control →

```
do_many_other_things();
```

```
}
```

# OpenMP synchronization: simple mutex locks

```
omp_lock_t lck;  
omp_init_lock(&lck);
```

```
#pragma omp parallel  
{
```

multiple threads of control →

```
do_many_things();
```

wait here for your turn ... →

```
omp_set_lock(&lck);
```

```
    cout << "here i am alone\n";
```

```
omp_unset_lock(&lck);
```

multiple threads of control →

```
do_many_other_things ();
```

```
}
```

```
omp_destroy_lock(&lck);
```

# requested number of threads: I really mean it ...

```
void main()
{
  int nthreads;
```

```
  omp_set_dynamic(0);
  omp_set_num_threads( omp_get_num_procs() );
  do_some_things();
```

NOTE: older versions use  
v omp\_num\_procs()

```
  #pragma omp parallel
  {
    int id = omp_get_thread_num();
    #pragma omp single
    {
      nthreads = omp_get_num_threads();
      if ( nthreads != what_you_want )
      {
        cout << "arrgh!" << endl;
        exit(fail);
      }
    }
    do_many_more_things();
  }
```

# OpenMP data environment

remember, young Jedi:

- global variables (meaning declared outside the scope of a parallel region) are **shared** among threads unless explicitly made private
- automatic variables declared within parallel region scope are **private**
- stack variable declared in functions called from within a parallel region are **private**

# data sharing attributes

`#pragma omp parallel private(x)`

- each thread receives its own **uninitialized** variable x
- the variable x falls out-of-scope after the parallel region
- a global variable with the same name is unaffected (3.0 and later)

`#pragma omp parallel firstprivate(x)`

- x must be a global-scope variable
- each thread receives a **by-value copy** of x
- the local x's fall out-of-scope after the parallel region
- the base global variable with the same name is unaffected (3.0 and later)

`#pragma omp parallel lastprivate(x)`

- x must be a global-scope variable
- each thread receives a **by-value copy** of x
- the base global variable of the same name is set to the last value of x for the last thread to finish



# OpenMP sections

```
#pragma omp parallel  
{  
  ...
```

multiple threads of control →  
each section assigned to a  
different thread

```
#pragma omp sections  
{  
  #pragma omp section  
    X_calculation();  
  #pragma omp section  
    y_calculation();  
  #pragma omp section  
    z_calculation();  
}  
...
```

with *nowait*:  
extra threads continue ahead

by default:  
extra threads are idled

```
}
```

# OpenMP schedule clause

The schedule clause determines how loop iterators are mapped onto threads:

`#pragma omp parallel for schedule( static [, chunk] )`

- chunking computed at compile time
- fixed-sized chunks assigned (alternating) to num\_threads
- typical default is:  $\text{chunk} = \text{iterations} / \text{num\_threads}$
- set `chunk = 1` for cyclic distribution

`#pragma omp parallel for schedule( dynamic [, chunk] )`

- run-time scheduling (with associated overhead)
- each thread grabs “chunk” iterations off queue until all iterations have been scheduled
- good load-balancing for uneven workloads

# OpenMP schedule clause

The schedule clause determines how loop iterators are mapped onto threads:

`#pragma omp parallel for schedule( guided[, chunk] )`

- threads dynamically grab blocks of iterations
- chunk size starts relatively large, to get all threads busy with good amortization of overhead
- subsequently, chunk size is reduced to produce good workload balance
- high overhead, best for very uneven workload loops

`#pragma omp parallel for schedule( runtime )`

- schedule and chunk size taken from environment variable or from runtime library routines

# OpenMP key env variables

unix environment variables of note:

OMP\_NUM\_THREADS

- set to the desired default number of threads
- **still just a request** and system can override

OMP\_SCHEDULE

- “schedule [, chunk]”

# cse5441 - parallel computing

## Open MP