


cse5441 - parallel computing



vectorizing compilers



not all compilers are created equally ...

array programming

array programming languages (such as Fortran 90) provide operators which generalize scalar functions to higher dimensions


Fortran 77

```
do i=1,n
  do j=1,n
    C(j,i) = A(j,i) + B(j,i)
  enddo
enddo
```

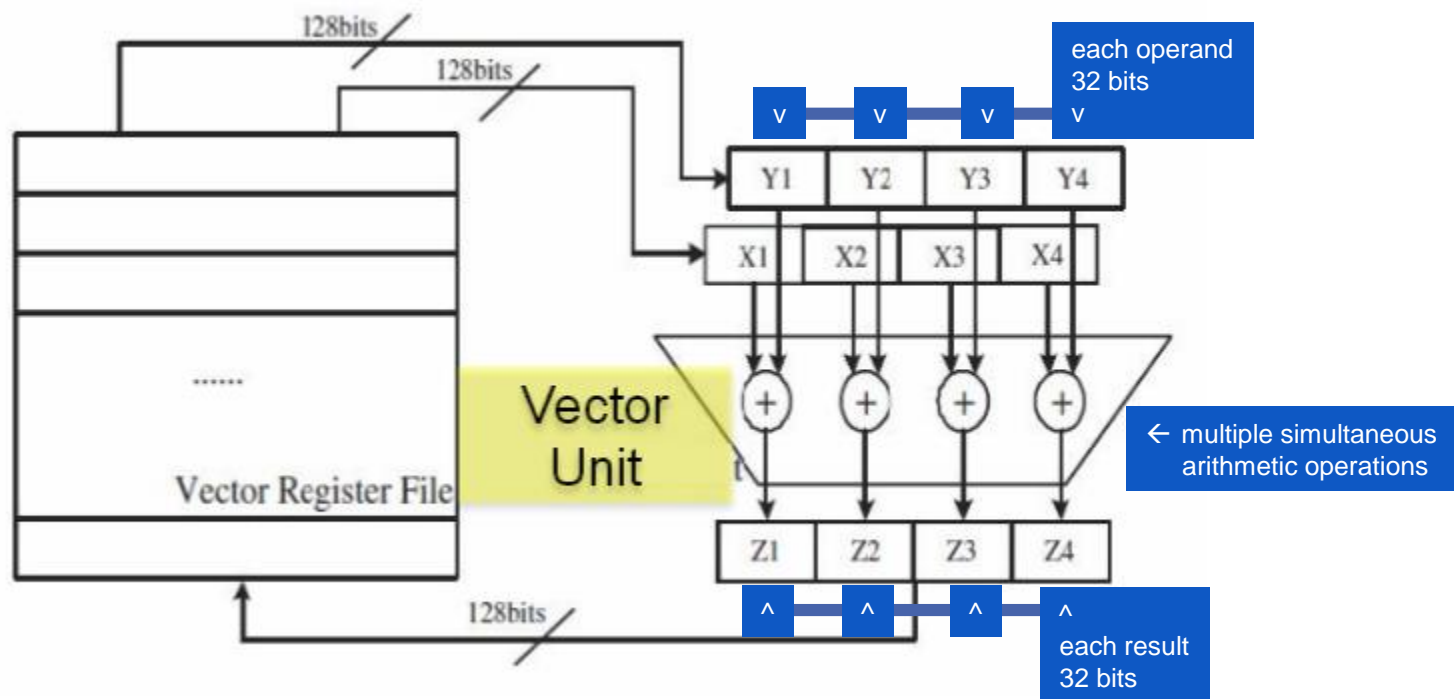
Fortran 90

```
C = A + B
```

scalar arithmetic / logic unit (ALU)



vector arithmetic / logic unit



unordered data distribution

```
for (iband = 0; iband < nbands; iband++)
{
    for (idir = 0; idir < ndir; idir++)
    {
        for (icell = 0; icell < ncells; icell++)
        {
            for (iface = 0; iface < nfcell[icell]; iface++)
            {
                if ( bface[curr] == 0 )           //interior face
                {
                    do_stuff(interior);
                }
                else                               //boundary face
                {
                    if ( bctype[ibface] == ADIA )
                    {
                        do_stuff(ADIA);
                    }
                    else if ( bctype[ibface] == ISOT )
                    {
                        do_stuff(ISOT);
                    }
                }
                }//end -- if interior or boundary case
                else
                . . .

            }//end -- loop over cell faces
        }//end -- cell loop
    }//end -- dir loop
}//end -- band loop
```

INPUT DATA (stylized)

INTER
INTER
ADIA
ISOT
PERM
INTER
PERM
ADIZ
ADIZ
INTER
PERM
ISOT
PERM
ADIA
ISOT
XPR
ISOT
PERM
ADIA
INTER

assume all inputs
are of same type,
and in these
application
categories

SIMD adaptation

UNORDERED INPUT DATA

INTER
INTER
ADIA
ISOT
PERM
INTER
PERM
ADIZ
ADIZ
INTER
PERM
ISOT
PERM
ADIA
ISOT
XPR
ISOT
PERM
ADIA
INTER

.
. .

PARTITIONED INPUT DATA

ADIA
ADIA
ADIA
ADIA
ADIA
ADIA
ADIA
ADIA
ADIA
ADIA

.
. .
INTER
INTER
INTER
INTER
INTER

.
. .

SIMD adaptation

```
for (iband = 0; iband < nbands; iband++)
{
    for (idir = 0; idir < ndir; idir++)
    {
        for (icell = 0; icell < ncells; icell++)
        {
            for (iface = 0; iface < nfcell[icell]; iface++)
            {
                if ( bface[curr] == 0 )           //interior face
                {
                    do_stuff(interior);
                }
                else                               //boundary face
                {
                    if ( bctype[ibface] == ADIA )
                    {
                        do_stuff(ADIA);
                    }
                    else if ( bctype[ibface] == ISOT )
                    {
                        do_stuff(ISOT);
                    }
                }
            }
        }
    }
}
//end -- if interior or boundary case
//end -- loop over cell faces
//end -- cell loop
//end -- dir loop
//end -- band loop
```

```
for (iband = 0; iband < nbands; iband++)
{
    for (idir = 0; idir < ndir; idir++)
    {
        for (iface = 0; iface < nf_max; iface++)
        {
            //process interior cell faces
            for (indx = 0; indx < num_if_cells[iface]; indx++)
            {
                do_stuff(interior);
            }
            //process ISOT cell faces
            for (indx = ISOT_offset;
                indx < ISOT_offset+num_isot_cells[iface];
                indx++)
            {
                do_stuff(ISOT);
            }
            //process ADIA cell faces
            for (indx = ADIA_offset;
                indx < ADIA_offset+num_adia_cells[iface];
                indx++)
            {
                do_stuff(ADIA);
            }
        }
    }
}
//end -- face loop
//end -- dir loop
//end -- band loop
```

C/C++ vectorizing compilers

```
for (int i = 0; i < n; i++)  
    c[i] = a[i] + b[i];
```

compilers automatically handle
the simple cases

```
for (int i = 0; i < n; i++)  
{  
    sum = 0.0;  
    for (int j = 0; j < n; j++)  
    {  
        sum += A[j][i];  
    }  
    B[i] = sum;  
}
```

what makes this loop more challenging?

- no stride-1 access
- sum creates a loop-carried dependency for i

C/C++ vectorizing compilers

example

```
for (int i = 0; i < n; i++)
{
    sum = 0.0;
    for (int j = 0; j < n; j++)
    {
        sum += A[j][i];
    }
    B[i] = sum;
}
```

A

```
for (int i = 0; i < n; i++)
{
    sum[i] = 0.0;
    for (int j = 0; j < n; j++)
    {
        sum[i] += A[j][i];
    }
    B[i] = sum[i];
}
```

B

scalar expansion:

eliminates loop dependency

```
for (int i = 0; i < n; i++)
{
    sum[i] = 0.0;
}

for (int j = 0; j < n; j++)
{
    for (int i = 0; i < n; i++)
    {
        sum[i] += A[j][i];
    }
}

for (int i = 0; i < n; i++)
{
    B[i] = sum[i];
}
```

C

plus
loop reordering
and distribution:

provides
stride-1
access

C/C++ vectorizing compilers

example

```
for (int i = 0; i < n; i++)
{
    sum = 0.0;
    for (int j = 0; j < n; j++)
    {
        sum += A[j][i];
    }
    B[i] = sum;
}
```

A

Intel Nehalem

loop not vectorized

IBM Power 7

loop not vectorized

```
for (int i = 0; i < n; i++)
{
    sum[i] = 0.0;
    for (int j = 0; j < n; j++)
    {
        sum[i] += A[j][i];
    }
    B[i] = sum[i];
}
```

B

Intel Nehalem

loop vectorized

speedup: 2.6 (62% faster)

relative run-time 0.6

IBM Power 7

loop interchanged
and vectorized

speedup: 2.0 (50% faster)


relative run-time 0.2

stripmine

step 1 of 2

ORIGINAL

```
for (int i = 0; i < n; i++)  
{  
S1 A[i] = B[i] + 1.0;  
S2 C[i] = A[i] + 2.0;  
}
```



STRIPMINE


```
for (int i = 0; i < n; i += stripsize)  
{  
    for (int j = i; j < i + stripsize; j++)  
    {  
        A[j] = B[j] + 1.0;  
        C[j] = A[j] + 2.0;  
    }  
}
```

stripmine - distribute

step 2 of 2


ORIGINAL

```
for (int i = 0; i < n; i++)  
{  
  S1 A[i] = B[i] + 1.0;  
  S2 C[i] = A[i] + 2.0;  
}
```



STRIPMINE

```
for (int i = 0; i < n; i += stripsize)  
{  
  for (int j = i; j < i + stripsize; j++)  
  {  
    A[j] = B[j] + 1.0;  
    C[j] = A[j] + 2.0;  
  }  
}
```



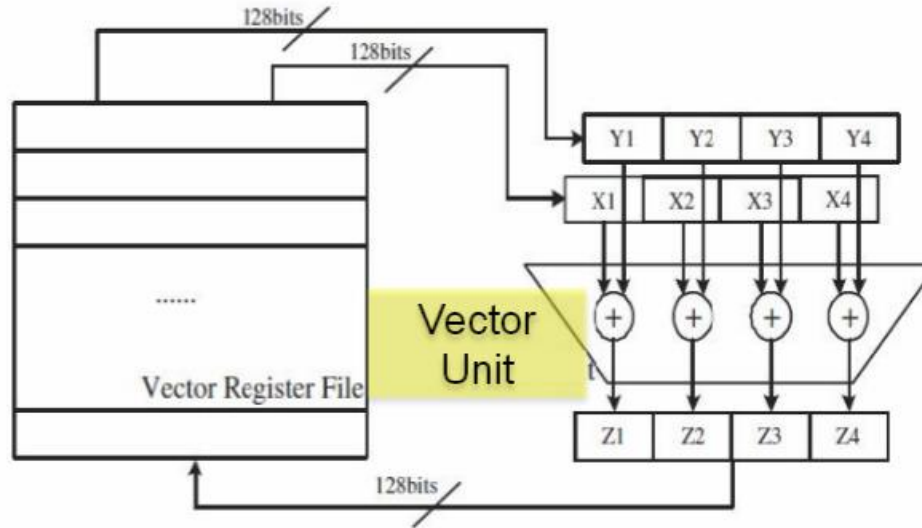
DISTRIBUTE

```
for (int i = 0; i < n; i += stripsize)  
{  
  for (int j = i; j < i + stripsize; j++)  
  {  
    A[j] = B[j] + 1.0;  
  }  
  for (int j = i; j < i + stripsize; j++)  
  {  
    C[j] = A[j] + 2.0;  
  }  
}
```

VECTORIZED

```
for (int i = k; i < n; i += q)  
{  
  A[i : i + q - 1] = B[i : i + q - 1] + 1.0;  
}  
for (int i = k; i < n; i += q)  
{  
  C[i : i + q - 1] = A[i : i + q - 1] + 2.0;  
}  
}
```

vectorization result



scalar:

```
load  r1  &operand1
load  r2  &operand2
add   r3  r1, r2
store r3  &result
load  r1  &operand1a
load  r2  &operand2a
add   r3  r1, r2
store r3  &result
```

scalar:

```
load  r1  &operand1b
load  r2  &operand2b
add   r3  r1, r2
store r3  &result
load  r1  &operand1c
load  r2  &operand2c
add   r3  r1, r2
store r3  &result
```


vector:

```
loadv  vr1  &operand1
loadv  vr2  &operand2
addv   vr3  vr1, vr2
storev vr3  &result
```

how well do compilers vectorize?

Compiler	XLC	ICC	GCC
Loops			
Total	159		
Vectorized	74	75	32
Not vectorized	85	84	127
Average Speed Up	1.73	1.85	1.30

Compiler	XLC but not ICC	ICC but not XLC
Loops		
Vectorized	25	26





adding manual vectorization hints increased the average speedup (IBM) from 1.73 to 3.78

acyclic dependence graphs

forward dependencies

```
for (int i = 0; i < max; i++)  
{  
  S1 a[ i ] = b[ i ] + c[ i ];  
  S2 d[ i ] = a[ i ] + 1;  
}
```




can we group all the $S1_n$
and follow with $S2_n$?

forward dependencies

example

```
for (int i = 0; i < max; i++)  
{  
  S1 a[ i ] = b[ i ] + c[ i ];  
  S2 d[ i ] = a[ i ] + 1;  
}
```





XLC: vectorized, speedup 2.0
ICC: vectorized, speedup 1.6

acyclic dependence graphs

backward dependencies

```
for (int i = 0; i < max; i++)  
{  
  S1 a[ i ] = b[ i ] + c[ i ];  
  S2 d[ i ] = a[ i+1 ] + 1;  
}
```





can we group all the $S2_n$
and follow with $S1_n$?

acyclic dependence graphs

backward dependencies

```
for (int i = 0; i < max; i++)  
{  
  S2 d[ i ] = a[ i+1 ] + 1;  
  S1 a[ i ] = b[ i ] + c[ i ];  
}
```



can we group all the $S2_n$
and follow with $S1_n$?

backward dependencies

example

original

```
for (int i = 0; i < max; i++)  
{  
    a[ i ] = b[ i ] + c[ i ];  
    d[ i ] = a[ i+1 ] + 1;  
}
```

re-ordered

```
for (int i = 0; i < max; i++)  
{  
    d[ i ] = a[ i+1 ] + 1;  
    a[ i ] = b[ i ] + c[ i ];  
}
```


ICC time 12.6
non-vectorized

XLC time 0.6
vectorized

ICC time 9.4
vectorized


XLC time 0.6
vectorized

cyclic dependence graphs




```
for (int i = 0; i < max; i++)  
{  
    a[ i ] = a[ i+1 ] + b[ i ];  
}
```

ICC
vectorized
XLC
vectorized



```
for (int i = 0; i < max; i++)  
{  
    a[ i ] = a[ i-1 ] + b[ i ];  
}
```

ICC
non-vectorized
XLC
non-vectorized



```
for (int i = 0; i < max; i++)  
{  
    a[ i ] = a[ i-4 ] + b[ i ];  
}
```

ICC
vectorized
XLC
vectorized

more on vectorizing compilers ...

references:

"Intel® Cilk™ Plus Support," [online] [_software.intel.com/en-us/articles/intel-cilk-plus-support](http://software.intel.com/en-us/articles/intel-cilk-plus-support)

"A Guide to Auto-vectorization with Intel® C++ Compilers," [online]
software.intel.com/en-us/articles/a-guide-to-auto-vectorization-with-intel-c-compilers

"Quick-Reference Guide to Optimization with Intel® Compilers version 13," [online]
software.intel.com/sites/default/files/Compiler_QRG_2013.pdf

"Intel® C++ Intrinsics Reference," [online]
software.intel.com/sites/products/documentation/studio/composer/en-us/2011/compiler_c/intref_cls/common/intref_bk_intro.htm

"Intel Advanced Vector Extensions (AVX)," [online] software.intel.com/en-us/avx

"Compiler Prefetching for the Intel® Xeon Phi™ coprocessor," [online]
software.intel.com/sites/default/files/managed/5d/f3/5.3-prefetching-on-mic-4.pdf

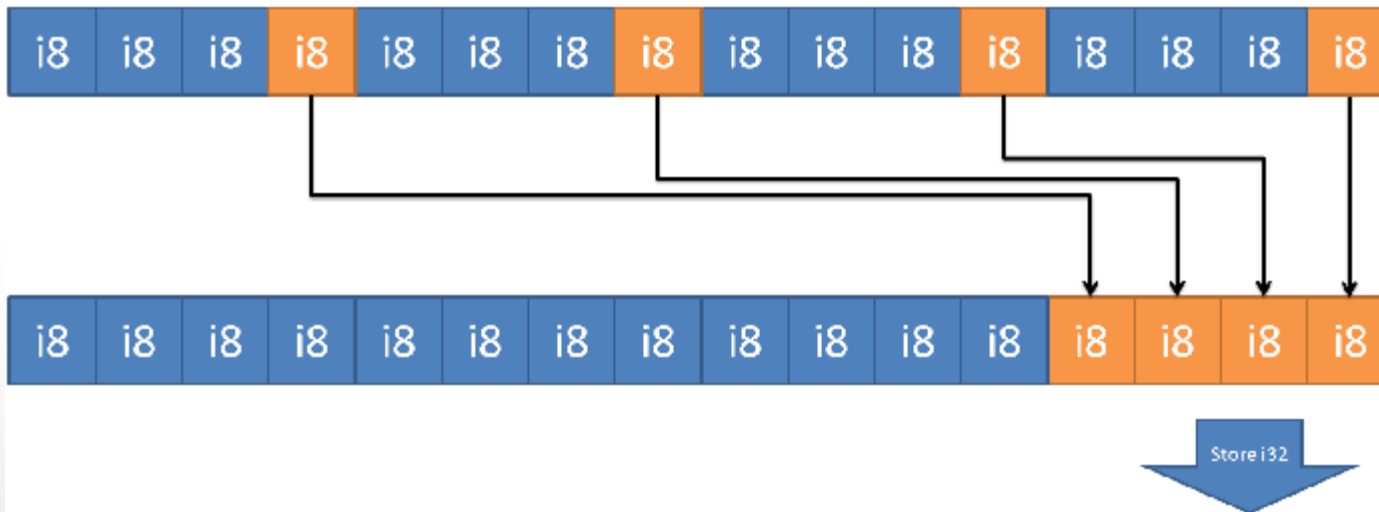
trade-offs



vectorization is not always profitable ...

cse5441 - parallel computing

vectorizing compilers



not all compilers are created equally ...